

Final Report

due November 16, 2021 by 11:59 PM

Enzo, Layla, Madeleine

November 16, 2021

```
cities <- initial_data %>%
  select(StateAbbr, PlaceName, Population2010, PlaceFIPS, Geolocation,
    ACCESS2_AdjPrev, CANCER_AdjPrev, CHD_AdjPrev, CHECKUP_AdjPrev,
    COPD_AdjPrev, COLON_SCREEN_AdjPrev, COREM_AdjPrev, COREW_AdjPrev,
    KIDNEY_AdjPrev, MAMMOUSE_AdjPrev, PAPTEST_AdjPrev) %>%
  dplyr::rename(state = StateAbbr, city = PlaceName, population = Population2010,
    health_access = ACCESS2_AdjPrev, cancer = CANCER_AdjPrev,
    heart_disease = CHD_AdjPrev, checkup = CHECKUP_AdjPrev,
    chronic_lung_disease = COPD_AdjPrev,
    colon_screen = COLON_SCREEN_AdjPrev,
    men_colorectal_cancer_screen = COREM_AdjPrev,
    women_colorectal_cancer_screen = COREW_AdjPrev,
    chronic_kidney_disease = KIDNEY_AdjPrev,
    mammogram = MAMMOUSE_AdjPrev,
    pap_test = PAPTEST_AdjPrev) %>%
  mutate(large_metro = (population >= 1500000),
    metro = (population >= 500000 & population < 1500000),
    med_urban = (population >= 200000 & population < 500000),
    small_urban = (population >= 50000 & population < 200000)) %>%
  mutate(west = (state %in% c("WA", "OR", "ID", "MT", "WY", "CA", "NV", "UT",
    "CO", "AZ", "NM")),
    midwest = (state %in% c("ND", "SD", "NE", "KS", "MN", "IA", "MO",
    "WI", "IL", "IN", "MI", "OH")),
    northeast = (state %in% c("PA", "NY", "VT", "NH", "MA", "CT", "RI",
    "NJ", "ME", "DC")),
    south = (state %in% c("OK", "TX", "AR", "LA", "MS", "AL", "TN", "KY",
    "WV", "VA", "MD", "DE", "NC", "SC", "GA",
    "FL"))) %>%
  mutate(region = ifelse(west == TRUE, "West",
    ifelse(midwest == TRUE, "Midwest",
    ifelse(northeast == TRUE, "Northeast",
    ifelse(south == TRUE,
    "South", NA)))) %>%
  mutate(city_size = ifelse(large_metro == TRUE, "Large Metropolitan",
    ifelse(metro == TRUE, "Metropolitan",
    ifelse(med_urban == TRUE, "Medium-Size Urban",
    ifelse(small_urban == TRUE,
    "Small-Size Urban", NA)))) %>%
  na.omit(city_size) %>%
  na.omit(region) %>%
  mutate(checkup_n = (population*(checkup/100)),
```

```

    colon_n_screened = (population*(colon_screen/100)),
    m_colorectal_n_screened = (population*(men_colorectal_cancer_screen/100)),
    w_colorectal_n_screened =
      (population*(women_colorectal_cancer_screen/100)),
    mammogram_n_screened = (population*(mammogram/100)),
    pap_n_screened = (population*(pap_test/100))) %>%
mutate(no_checkup_n = (population - checkup_n),
       no_colon_n_screened = (population - colon_n_screened),
       no_m_colorectal_n_screened = (population - m_colorectal_n_screened),
       no_w_colorectal_n_screened =
         (population - w_colorectal_n_screened),
       no_mammogram_n_screened = (population - mammogram_n_screened),
       no_pap_n_screened = (population - pap_n_screened))

# checked the residual plot without this and need to do this to make model more accurate
#this is the training data set called model_cities
set.seed(100)

split <- initial_split(cities, prop = 3/4, strata = region)

model_cities <- training(split)
cities_test <- testing(split)

#for introductory visualizations
west_cities <- cities %>%
  filter(west)

midwest_cities <- cities %>%
  filter(midwest)

northeast_cities <- cities %>%
  filter(northeast)

south_cities <- cities %>%
  filter(south)

#boxplot of region and city size

```


Introductory Visualisations of Data Set

```

cities2 <- cities %>%
  filter( region != "NA") %>%
  mutate(logcheckup_n = log(checkup_n))

#boxplots of checkup (n and log) by region
ggplot(data = cities2,
       aes(x = region, y = logcheckup_n,
           color = region)) +
  geom_boxplot()

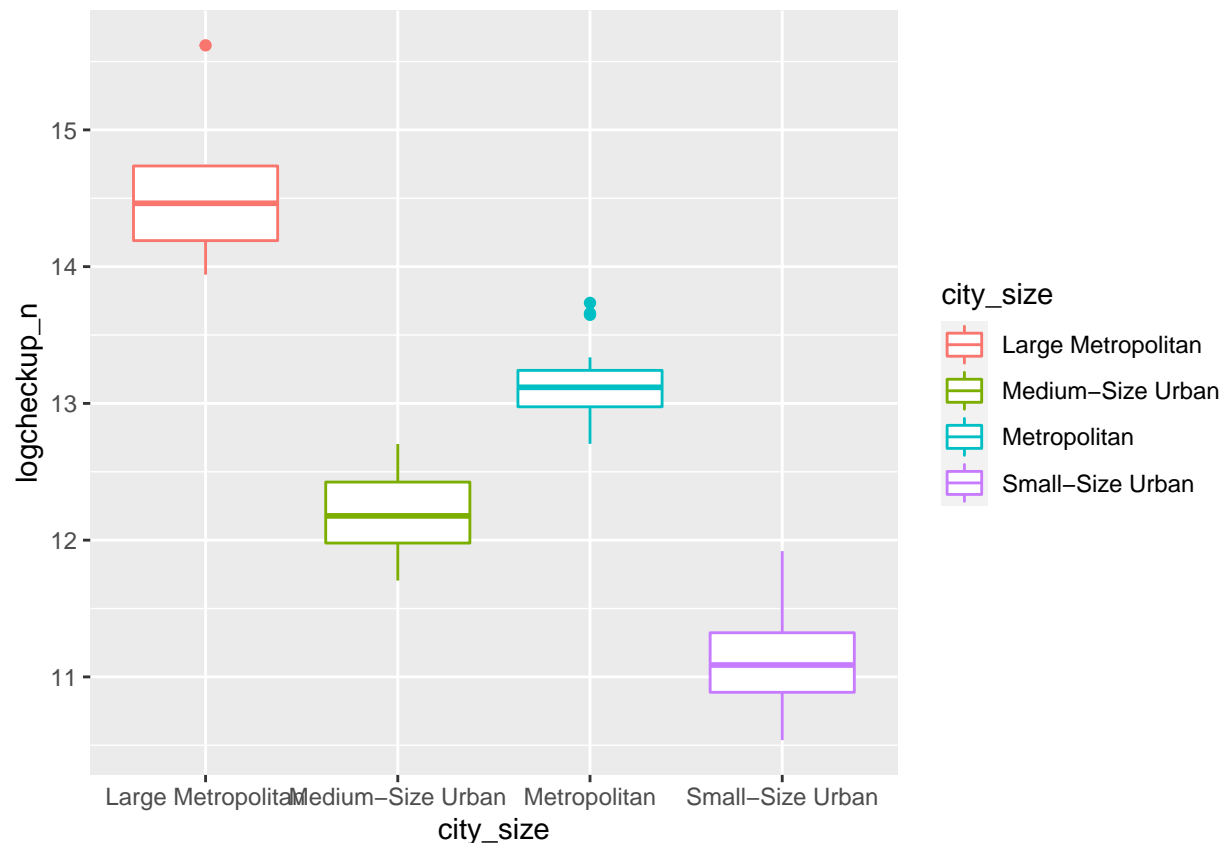
```



```
#anova shows means of checkup (n and log) are different among the regions
anova1 <- aov(logcheckup_n ~ region, data = cities2)
tidy(anova1)
```

```
## # A tibble: 2 x 6
##   term      df  sumsq meansq statistic  p.value
##   <chr>    <dbl> <dbl>  <dbl>    <dbl>    <dbl>
## 1 region      3   6.32   2.11     4.16  0.00638
## 2 Residuals 447 226.    0.507     NA     NA
```

```
#boxplots of checkup (n and log) by city size
ggplot(data = cities2,
       aes(x = city_size, y = logcheckup_n,
           color = city_size)) +
  geom_boxplot()
```




```
#anova shows means of checkup (n and log) are different among the city sizes
anova2 <- aov(logcheckup_n ~ city_size, data = cities2)
tidy(anova2)
```

```
## # A tibble: 2 x 6
##   term      df sumsq meansq statistic    p.value
##   <chr>   <dbl> <dbl>   <dbl>     <dbl>    <dbl>
## 1 city_size     3 191.   63.5      671. 4.15e-165
## 2 Residuals   447  42.3    0.0946     NA      NA
```


```
#enzo's graph for health access vs checkup
ggplot(data = cities2,
       aes(x = health_access, y = logcheckup_n)) +
  geom_col(aes(width = 1))
```


```
## Warning: Ignoring unknown aesthetics: width
```

```
## Warning: position_stack requires non-overlapping x intervals
```




```
#madeleine's graph for health access vs checkup  
ggplot(cities, aes(x = health_access,  
                   y = checkup))+  
  geom_point()
```



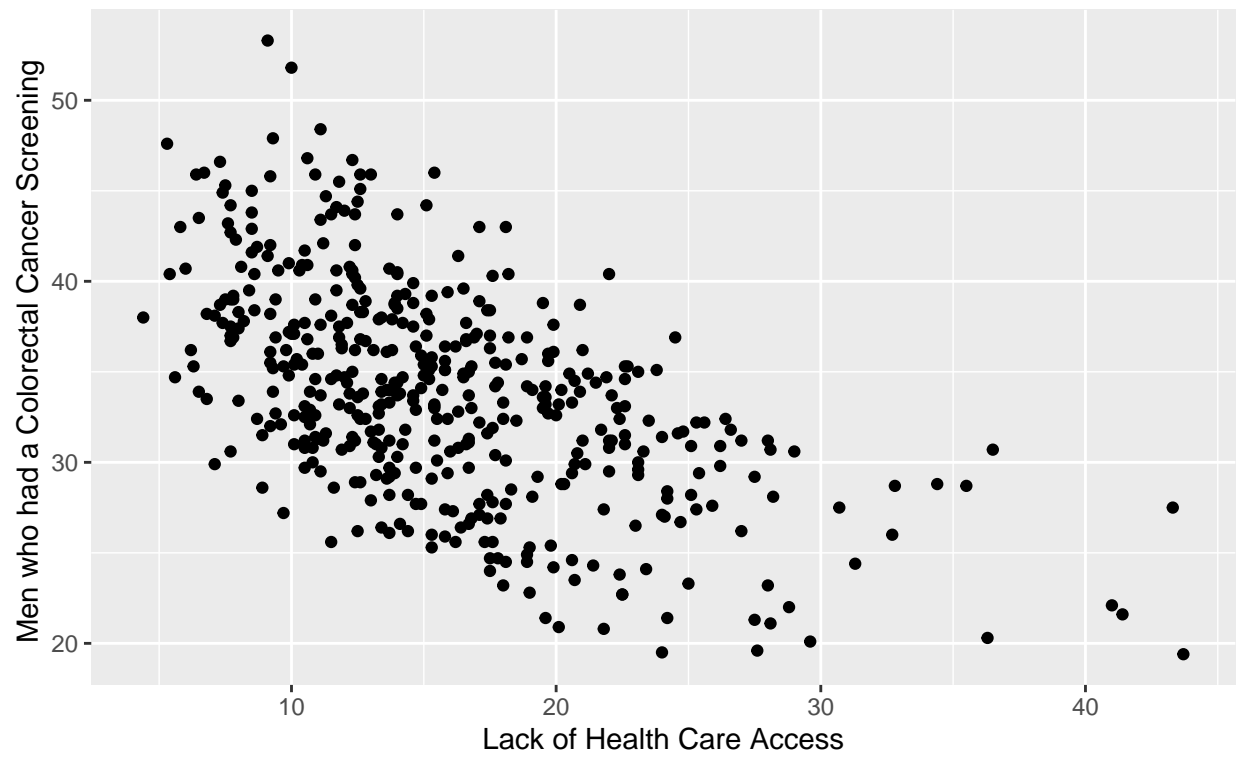


Number of Checkups by City Size




Colorectal Cancer Screening and Lack of Healthcare Access

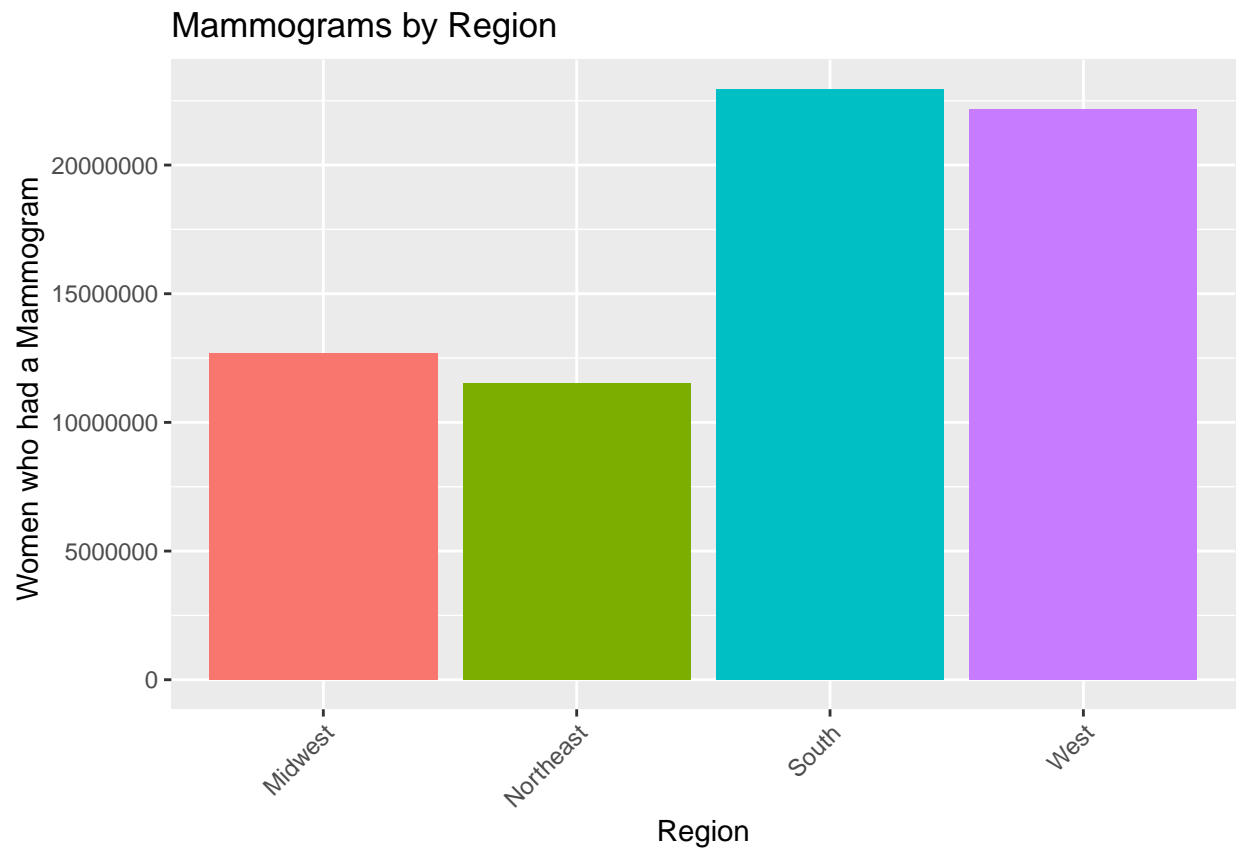
Men only



Colorectal Cancer Screening and Lack of Healthcare Access

Women only






```
cities2 <- cities %>%
  filter( region != "NA")

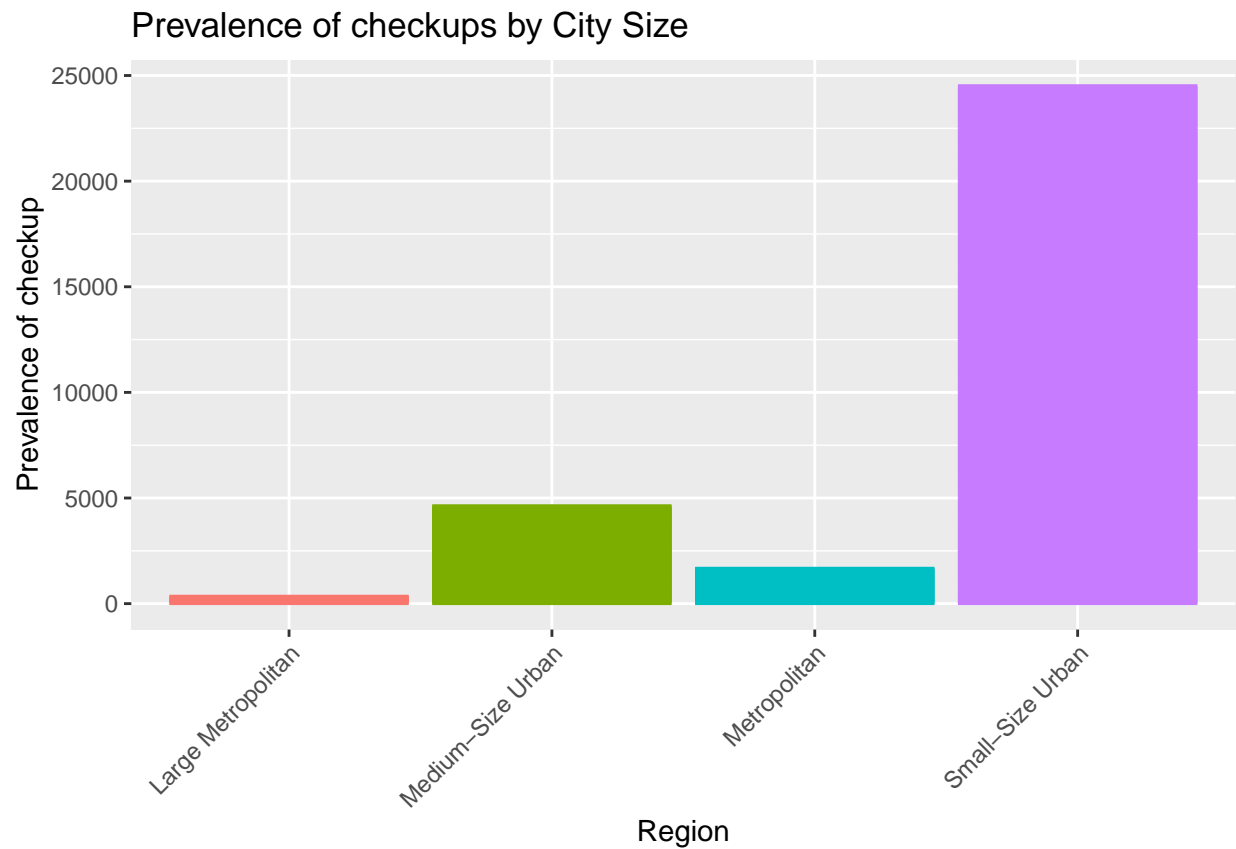
ggplot(cities2, aes(x = region,
                    y = checkup,
                    color = region))+
  geom_bar( stat = 'identity' )+
  labs( x = "Region",
        y = "Prevalence of checkup",
        title = "Prevalence of checkups by Region") +
  theme(axis.text.x = element_text(angle = 45,hjust=1))+
  theme(legend.position = "none")
```

Prevalence of checkups by Region



```
cities2 <- cities %>%
  filter( city_size != "NA")


ggplot(cities2, aes(x = city_size,
                    y = checkup,
                    color = city_size))+
  geom_bar( stat = 'identity' )+
  labs( x = "Region",
        y = "Prevalence of checkup",
        title = "Prevalence of checkups by City Size") +
  theme(axis.text.x = element_text(angle = 45,hjust=1))+
  theme(legend.position = "none")
```



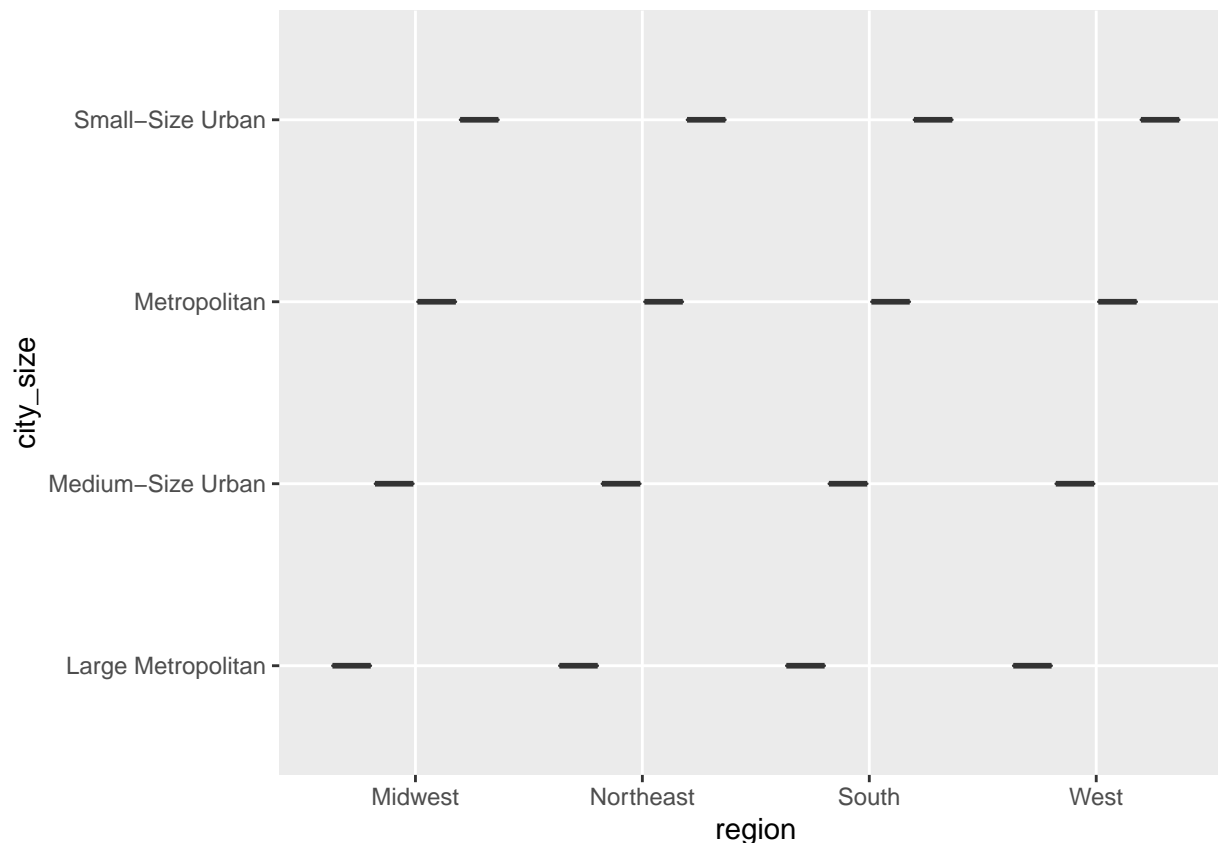
#doesn't look that helpful tbh won't hurt my feelings if u trash it

```
ggplot(cities, aes(x = health_access,  
                   y = checkup))+  
  geom_point()+  
  labs( x = "Lack of Health Care Access",  
        y = "Prevalence of Checkups",  
        title = "Prevalence of Checkups and Lack of Healthcare Access")
```

Prevalence of Checkups and Lack of Healthcare Access



#cloud of points means something-- don't forget this later on

[illegible]

```
#head(cities_heatmap)
# ggplot(cities, aes(checkup, )) +
```

```
correlation1 <- cor.test(cities$checkup_n, cities$pap_n_screened,  
                        method = "pearson")  
correlation2 <- cor.test(cities$checkup_n, cities$colon_n_screened,  
                        method = "pearson")  
correlation3 <- cor.test(cities$checkup_n, cities$mammogram_n_screened,  
                        method = "pearson")  
print(correlation1)
```

```
##
## Pearson's product-moment correlation
##
## data: cities$checkup_n and cities$pap_n_screened
## t = 332.26, df = 449, p-value < 0.00000000000000022
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.9975606 0.9983151
## sample estimates:
## cor
## 0.9979726
```

```
print(correlation2)
```

```
##
## Pearson's product-moment correlation
##
## data: cities$checkup_n and cities$colon_n_screened
## t = 294.91, df = 449, p-value < 0.00000000000000022
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.9969063 0.9978629
## sample estimates:
## cor
## 0.9974287
```

```
print(correlation3)
```

```
##
## Pearson's product-moment correlation
##
## data: cities$checkup_n and cities$mammogram_n_screened
## t = 437.63, df = 449, p-value < 0.00000000000000022
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.9985920 0.9990276
## sample estimates:
## cor
## 0.9988299
```

Multiple Logistic Regression Model

```
##
## Call:
## glm(formula = cbind(checkup_n, no_checkup_n) ~ region + city_size +
## health_access, family = binomial, data = model_cities)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -99.359  -18.975   -0.627   14.963  130.327
##
## Coefficients:
##              Estimate Std. Error z value
## (Intercept)    0.97037012  0.00122436  792.555
## regionNortheast    0.19874789  0.00093326  212.960
## regionSouth       0.10411063  0.00084367  123.402
## regionWest       -0.27722312  0.00071380 -388.374
## city_sizeMedium-Size Urban  0.01336343  0.00086185  15.506
## city_sizeMetropolitan -0.00289017  0.00084854   -3.406
## city_sizeSmall-Size Urban -0.01782009  0.00077588 -22.968
## health_access    -0.00596421  0.00005375 -110.965
##
##              Pr(>|z|)
## (Intercept) < 0.0000000000000002 ***
## regionNortheast < 0.0000000000000002 ***
## regionSouth < 0.0000000000000002 ***
## regionWest < 0.0000000000000002 ***
```



```

## city_sizeMedium-Size Urban < 0.0000000000000002 ***
## city_sizeMetropolitan          0.000659 ***
## city_sizeSmall-Size Urban < 0.0000000000000002 ***
## health_access < 0.0000000000000002 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 808649  on 336  degrees of freedom
## Residual deviance: 290529  on 329  degrees of freedom
## AIC: 294628
##
## Number of Fisher Scoring iterations: 3
##
## Call:
## glm(formula = cbind(checkup_n, no_checkup_n) ~ region + city_size +
##     health_access + region * city_size, family = binomial, data = model_cities)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -103.564   -17.554    0.418    14.078    86.320
##
## Coefficients:
##              Estimate Std. Error z value
## (Intercept)      0.97009830  0.00162902  595.512
## regionNortheast    0.15804724  0.00156374  101.070
## regionSouth      -0.01826191  0.00211966   -8.616
## regionWest       -0.21440752  0.00172584 -124.234
## city_sizeMedium-Size Urban -0.05649839  0.00178519  -31.648
## city_sizeMetropolitan  0.09483272  0.00188009   50.440
## city_sizeSmall-Size Urban -0.06367042  0.00163329  -38.983
## health_access    -0.00488135  0.00005675  -86.017
## regionNortheast:city_sizeMedium-Size Urban  0.16894949  0.00322129   52.448
## regionSouth:city_sizeMedium-Size Urban  0.20281749  0.00254564   79.672
## regionWest:city_sizeMedium-Size Urban -0.01553915  0.00228157   -6.811
## regionNortheast:city_sizeMetropolitan  0.10447464  0.00301573   34.643
## regionSouth:city_sizeMetropolitan -0.04324966  0.00254176  -17.016
## regionWest:city_sizeMetropolitan -0.21259142  0.00239770  -88.665
## regionNortheast:city_sizeSmall-Size Urban  0.04923987  0.00237171   20.761
## regionSouth:city_sizeSmall-Size Urban  0.18529679  0.00234383   79.057
## regionWest:city_sizeSmall-Size Urban -0.04092672  0.00205398  -19.926
##
##              Pr(>|z|)
## (Intercept) < 0.0000000000000002 ***
## regionNortheast < 0.0000000000000002 ***
## regionSouth < 0.0000000000000002 ***
## regionWest < 0.0000000000000002 ***
## city_sizeMedium-Size Urban < 0.0000000000000002 ***
## city_sizeMetropolitan < 0.0000000000000002 ***
## city_sizeSmall-Size Urban < 0.0000000000000002 ***
## health_access < 0.0000000000000002 ***
## regionNortheast:city_sizeMedium-Size Urban < 0.0000000000000002 ***
## regionSouth:city_sizeMedium-Size Urban < 0.0000000000000002 ***

```

```
## regionWest:city_sizeMedium-Size Urban      0.000000000000971 ***
## regionNortheast:city_sizeMetropolitan      < 0.0000000000000002 ***
## regionSouth:city_sizeMetropolitan          < 0.0000000000000002 ***
## regionWest:city_sizeMetropolitan           < 0.0000000000000002 ***
## regionNortheast:city_sizeSmall-Size Urban < 0.0000000000000002 ***
## regionSouth:city_sizeSmall-Size Urban     < 0.0000000000000002 ***
## regionWest:city_sizeSmall-Size Urban      < 0.0000000000000002 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 808649  on 336  degrees of freedom
## Residual deviance: 252886  on 320  degrees of freedom
## AIC: 257003
##
## Number of Fisher Scoring iterations: 3
```

```
checkup_fit <- logistic_reg() %>%
```

```
set_engine("glm") %>%
```

```
fit(as.factor(checkup_n) ~ region + city_size + health_access,
data=cities,
```

```
family="binomial", exponentiate = TRUE)
```

```
result <- tidy(checkup_fit, conf.int=TRUE)
```

```
print(result)
```

#how do i visualize all the log reg models can i augment when using a log reg model that uses cbind? can i make a regression plot without using augment function?

```
##
## Call:
## glm(formula = cbind(colon_n_screened, no_colon_n_screened) ~
##      region + city_size + health_access, family = binomial, data = model_cities)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -98.034  -13.699   -1.068   18.543   65.912
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.96017646  0.00115479  831.47 <0.0000000000000002
## regionNortheast  0.01927332  0.00086022   22.41 <0.0000000000000002
## regionSouth     0.24458352  0.00079496  307.67 <0.0000000000000002
## regionWest      0.11740260  0.00068781  170.69 <0.0000000000000002
```

```

## city_sizeMedium-Size Urban 0.02151338 0.00081370 26.44 <0.0000000000000002
## city_sizeMetropolitan 0.02745954 0.00079974 34.34 <0.0000000000000002
## city_sizeSmall-Size Urban 0.02848220 0.00073139 38.94 <0.0000000000000002
## health_access -0.03364736 0.00005098 -660.05 <0.0000000000000002
##
## (Intercept) ***
## regionNortheast ***
## regionSouth ***
## regionWest ***
## city_sizeMedium-Size Urban ***
## city_sizeMetropolitan ***
## city_sizeSmall-Size Urban ***
## health_access ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 763964 on 336 degrees of freedom
## Residual deviance: 221991 on 329 degrees of freedom
## AIC: 226119
##
## Number of Fisher Scoring iterations: 3
##
## Call:
## glm(formula = cbind(mammogram_n_screened, no_mammogram_n_screened) ~
## region + city_size + health_access, family = binomial, data = model_cities)
##
## Deviance Residuals:
## Min 1Q Median 3Q Max
## -91.766 -17.040 4.115 18.677 90.943
##
## Coefficients:
## Estimate Std. Error z value Pr(>|z|)
## (Intercept) 1.2647938 0.0012665 998.664 < 0.0000000000000002
## regionNortheast 0.1026001 0.0009550 107.435 < 0.0000000000000002
## regionSouth 0.1759165 0.0008741 201.243 < 0.0000000000000002
## regionWest 0.0265853 0.0007513 35.384 < 0.0000000000000002
## city_sizeMedium-Size Urban -0.0133440 0.0008966 -14.883 < 0.0000000000000002
## city_sizeMetropolitan -0.0038478 0.0008831 -4.357 0.0000132
## city_sizeSmall-Size Urban -0.0075516 0.0008070 -9.357 < 0.0000000000000002
## health_access -0.0166378 0.0000557 -298.684 < 0.0000000000000002
##
## (Intercept) ***
## regionNortheast ***
## regionSouth ***
## regionWest ***
## city_sizeMedium-Size Urban ***
## city_sizeMetropolitan ***
## city_sizeSmall-Size Urban ***
## health_access ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 401462 on 336 degrees of freedom
## Residual deviance: 288696 on 329 degrees of freedom
## AIC: 292760
##
## Number of Fisher Scoring iterations: 3

##
## Call:
## glm(formula = cbind(pap_n_screened, no_pap_n_screened) ~ region +
## city_size + health_access, family = binomial, data = model_cities)
##
## Deviance Residuals:
## Min 1Q Median 3Q Max
## -153.489 -15.237 2.326 16.525 85.826
##
## Coefficients:
## Estimate Std. Error z value Pr(>|z|)
## (Intercept) 1.66586484 0.00133834 1244.73 <0.0000000000000002
## regionNortheast -0.07942884 0.00100571 -78.98 <0.0000000000000002
## regionSouth 0.12056658 0.00092792 129.93 <0.0000000000000002
## regionWest 0.01680779 0.00080678 20.83 <0.0000000000000002
## city_sizeMedium-Size Urban -0.10257377 0.00094715 -108.30 <0.0000000000000002
## city_sizeMetropolitan -0.04724339 0.00093719 -50.41 <0.0000000000000002
## city_sizeSmall-Size Urban -0.06835666 0.00085439 -80.01 <0.0000000000000002
## health_access -0.02159918 0.00005833 -370.29 <0.0000000000000002
##
## (Intercept) ***
## regionNortheast ***
## regionSouth ***
## regionWest ***
## city_sizeMedium-Size Urban ***
## city_sizeMetropolitan ***
## city_sizeSmall-Size Urban ***
## health_access ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 392788 on 336 degrees of freedom
## Residual deviance: 242902 on 329 degrees of freedom
## AIC: 246928
##
## Number of Fisher Scoring iterations: 3
```

Model Validation

model that was made from the training data set compared to the rest of the test data points from the cities data

```

library(pROC)

## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'
##
## The following objects are masked from 'package:stats':
##
##      cov, smooth, var
# predictions from no interaction model


test_prob1 = predict(checkup_fit, newdata = cities_test, type = "response")
print(test_prob1)

##      1      2      3      4      5      6      7      8
## 0.6455606 0.7314895 0.7011851 0.7228949 0.7112079 0.7050252 0.7078892 0.6422785
##      9     10     11     12     13     14     15     16
## 0.6523540 0.7276781 0.6500896 0.7057882 0.7093670 0.6459699 0.6957643 0.7177244
##     17     18     19     20     21     22     23     24
## 0.7083823 0.7351469 0.7208554 0.7154229 0.7243224 0.6459699 0.7294139 0.6503223
##     25     26     27     28     29     30     31     32
## 0.7157870 0.7352215 0.6481492 0.7284711 0.7143288 0.7110314 0.6446047 0.6527981
##     33     34     35     36     37     38     39     40
## 0.7096128 0.7054165 0.7094899 0.7451987 0.7226520 0.6461062 0.7220539 0.7156657
##     41     42     43     44     45     46     47     48
## 0.7207354 0.7451987 0.7174826 0.7044241 0.7430410 0.7104724 0.6389831 0.7072722
##     49     50     51     52     53     54     55     56
## 0.7123092 0.6467878 0.6435109 0.7165145 0.7284711 0.6454241 0.7445186 0.7242033
##     57     58     59     60     61     62     63     64
## 0.7335473 0.6511356 0.7117661 0.7243224 0.7476684 0.7315275 0.7064071 0.6512711
##     65     66     67     68     69     70     71     72
## 0.6370547 0.7013101 0.6485571 0.7217288 0.7072722 0.6484598 0.7081709 0.6439212
##     73     74     75     76     77     78     79     80
## 0.6297147 0.6896344 0.6297147 0.7254268 0.7297669 0.7100500 0.7293296 0.7290941
##     81     82     83     84     85     86     87     88
## 0.6609976 0.6507291 0.7171198 0.6437845 0.7104724 0.7486131 0.6499197 0.7189311
##     89     90     91     92     93     94     95     96
## 0.7278425 0.6500510 0.7045482 0.6411815 0.6967931 0.7279606 0.7207354 0.6476049
##     97     98     99    100    101    102    103    104
## 0.7131987 0.6388455 0.6420044 0.7485008 0.7246796 0.6462426 0.7173617 0.6443314
##    105    106    107    108    109    110    111    112
## 0.6432372 0.7183281 0.7304339 0.6462426 0.6435109 0.6482852 0.6330461 0.6476049
##    113    114
## 0.7081358 0.7275599

new_cities_test <-
  merge(cities_test, test_prob1, by = "row.names", all.x = TRUE)

ggplot(cities_test, aes(x = test_prob1, y = checkup)) +
  geom_point()

```



```
# prediction from model with interaction
```


```
test_prob2 = predict(checkup_fit2, newdata = cities_test, type = "response")
print(test_prob2)
```

```
##          1          2          3          4          5          6          7          8
## 0.6431587 0.7294199 0.6952775 0.7132379 0.7035892 0.7152382 0.7008332 0.6404655
##          9         10         11         12         13         14         15         16
## 0.6487404 0.7334590 0.6484848 0.6990904 0.7020599 0.6434947 0.7078243 0.7256013
##         17         18         19         20         21         22         23         24
## 0.7012424 0.7235076 0.7281210 0.7237509 0.7309143 0.6434947 0.7350225 0.6470701
##         25         26         27         28         29         30         31         32
## 0.7240436 0.7324917 0.6452845 0.7342612 0.7228717 0.7202237 0.6423741 0.6507070
##         33         34         35         36         37         38         39         40
## 0.7022641 0.6987823 0.7021620 0.7407306 0.7295681 0.6436067 0.7290863 0.7239460
##         41         42         43         44         45         46         47         48
## 0.7280244 0.7407306 0.7254069 0.6979597 0.7389454 0.7029780 0.6377634 0.7003212
##         49         50         51         52         53         54         55         56
## 0.7045046 0.6441664 0.6414765 0.7246284 0.7342612 0.6430467 0.7401677 0.7308183
##         57         58         59         60         61         62         63         64
## 0.7382032 0.6477387 0.7208134 0.7309143 0.7427469 0.7367302 0.6996036 0.6478500
##         65         66         67         68         69         70         71         72
## 0.6361832 0.6953809 0.6456196 0.7286605 0.7003212 0.6471484 0.6972072 0.6418132
##         73         74         75         76         77         78         79         80
## 0.6301744 0.7030993 0.6301744 0.7316419 0.7353076 0.7194361 0.7347929 0.7346026
##         81         82         83         84         85         86         87         88
```

```
## 0.6574440 0.6474045 0.7251151 0.6417010 0.7029780 0.7624053 0.6409192 0.7265722
##      89      90      91      92      93      94      95      96
## 0.7264229 0.6468471 0.6980626 0.6395658 0.6916457 0.7265199 0.7280244 0.6448374
##      97      98      99     100     101     102     103     104
## 0.7014155 0.6376507 0.6402407 0.7623169 0.7312022 0.6437187 0.7253097 0.6421498
##     105     106     107     108     109     110     111     112
## 0.6412520 0.7260870 0.7285519 0.6437187 0.6414765 0.6453962 0.6329005 0.6448374
##     113     114
## 0.7010378 0.7333636
```

```
new_cities_test <-
  merge(cities_test, test_prob2, by = "row.names", all.x = TRUE)

ggplot(cities_test, aes(x = test_prob2, y = checkout)) +
  geom_point()
```



```
#calculate root mean standard error of both models
```

```
rmse(new_cities_test, truth = checkout, estimate = test_prob1)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      69.0
```

```
rmse(new_cities_test, truth = checkout, estimate = test_prob2)
```

```
## # A tibble: 1 x 3
```

```
##      .metric .estimator .estimate
##      <chr>   <chr>         <dbl>
## 1 rmse      standard      69.0
```